



# Variational Autoencoders (VAEs)

Unlocking the power of probabilistic generative models in deep learning



## ◆ What is a VAE?

A **Variational Autoencoder (VAE)** is a [generative model](#) that learns to encode input data into a **probabilistic latent space** and then **decode** it back to reconstruct or generate new data.

It's based on deep learning and **probabilistic inference**, allowing the generation of new, similar data samples whilst maintaining the underlying data distribution.

- ❑  **In short:** VAEs learn *distributions* of data, not just compressed codes – enabling creativity and diversity in generation.



## ◆ Why Were VAEs Introduced?

### Traditional Autoencoders Can:

- Compress and reconstruct data
- Learn feature representations
- Reduce dimensionality

### But They **Cannot**:

- Generate truly *new* data
- Model the underlying **probability distribution**
- Sample meaningfully from latent space

💡 **VAE fixes this** by learning a **probability distribution** (mean  $\mu$  and variance  $\sigma$ ), allowing **sampling** from this learned distribution to generate *new, similar data*.



## ◆ Architecture of a VAE

A VAE consists of **three main components** working together:

### Encoder

Maps input data → parameters  
( $\mu, \sigma$ ) of a Gaussian distribution

**Analogy:** The [Analyst](#) who  
describes the face

### Latent Space

Compressed, abstract space of  
learned features

**Analogy:** The [Blueprint Room](#)  
storing "feature sketches"

### Decoder

Reconstructs or generates data from sampled latent vectors

**Analogy:** The [Painter](#) who creates a new face

■ **Key Idea:** Instead of outputting a single latent vector, the encoder  
outputs a [distribution](#) – giving the model flexibility and creativity.



# The Painting Studio Analogy

01

## Encoder Stage

Looks at a face and encodes it as probabilities ( $\mu, \sigma$ )

*Analyst describing features:* "eyes ~2±0.5 apart"

02

## Latent Space

Stores these probabilistic descriptions

*Blueprint room – similar faces are close together*

03

## Decoder Stage

Uses a sampled description to recreate or imagine a new face

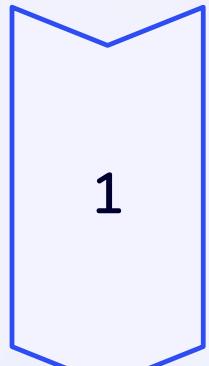
*Painter generating the full portrait*



**Result:** Even small tweaks in latent descriptions → slightly new, realistic faces with natural variations.

# ◆ Step-by-Step Example

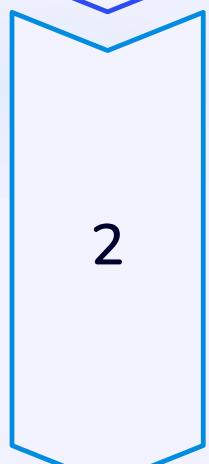
**Input:** A grayscale face image (e.g., 28×28 pixels)



## Step 1: Encode the Image

The encoder outputs:

- $\mu = [1.3, -0.7]$
- $\sigma = [0.2, 0.5]$



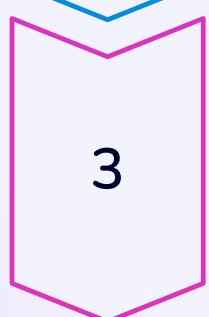
## Step 2: Sample from Latent Space

We draw a random vector:

$$z = \mu + \sigma \times \varepsilon, \text{ where } \varepsilon \sim N(0, 1)$$

Example:  $z = [1.4, -0.6]$

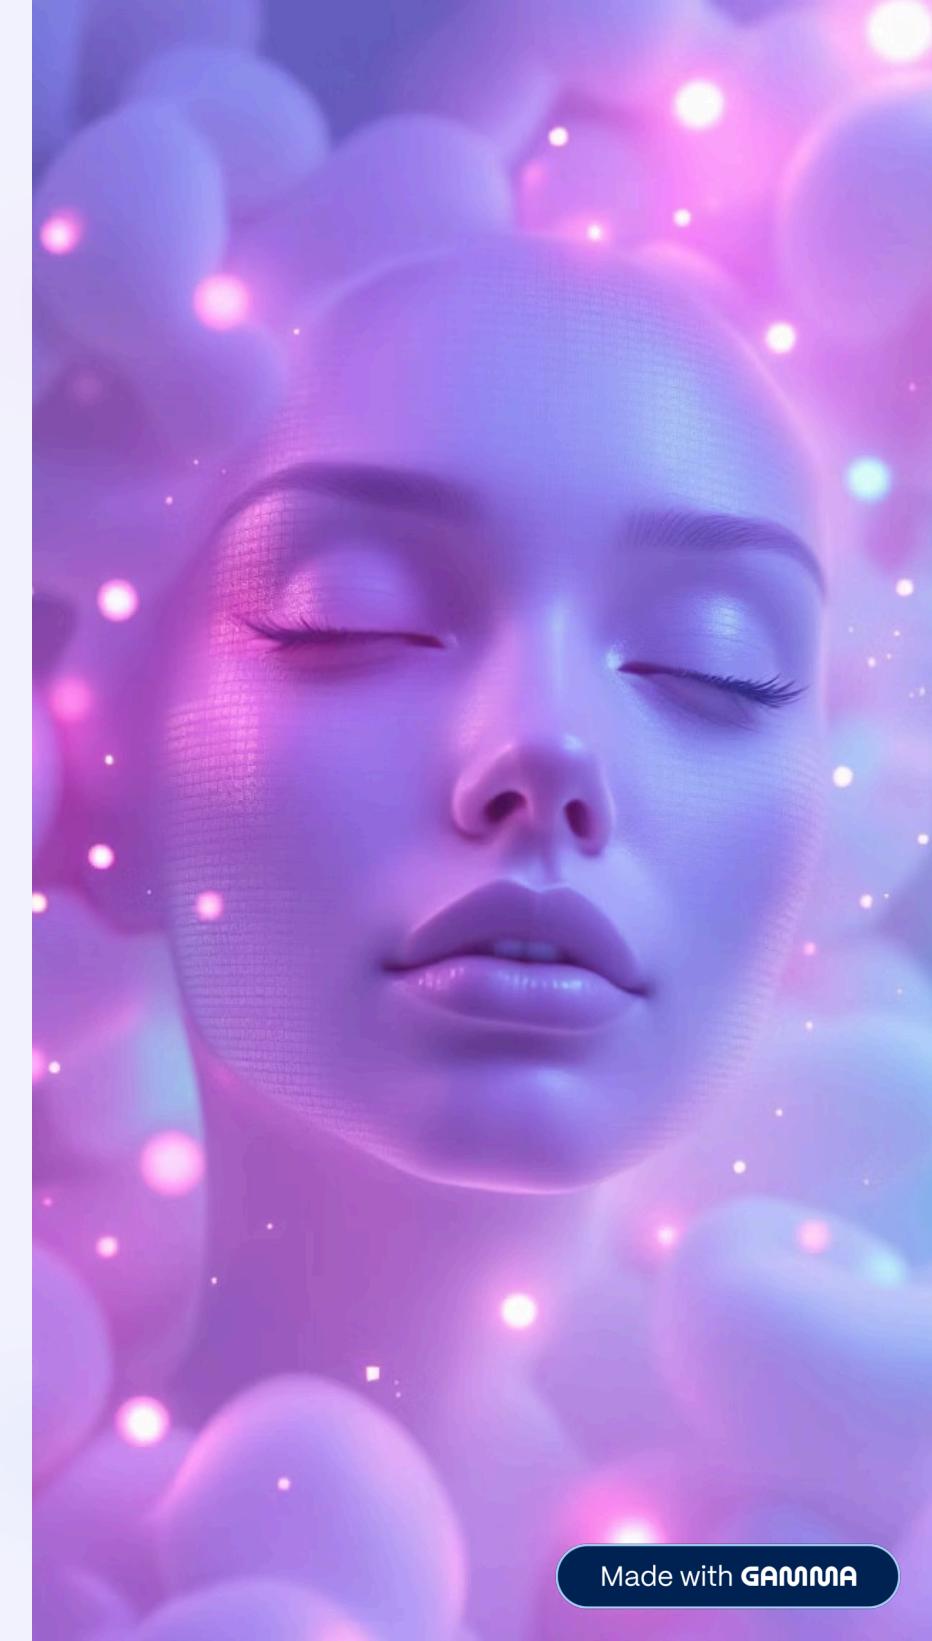
This randomness ensures **diverse** outputs.

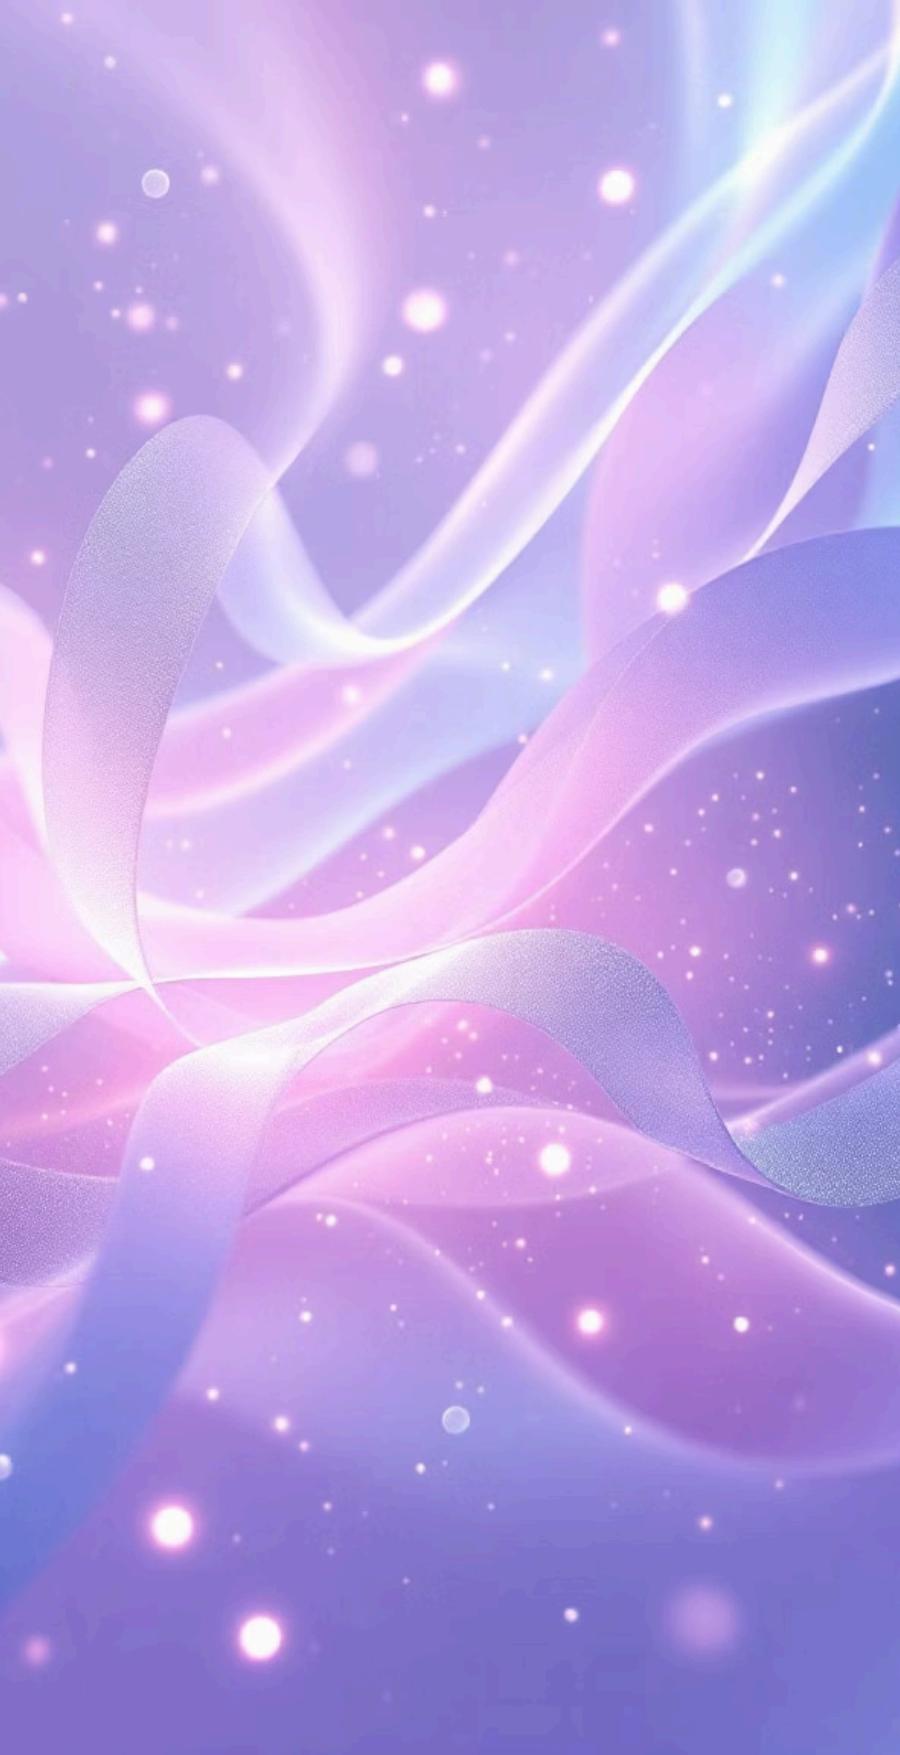


## Step 3: Decode the Sample

The decoder reconstructs an image from z

Ideally, it looks like the input (or a plausible new face)





# ◆ Why VAEs Are Powerful in Generative AI

## New Data Generation

Sample random  $z$  values → produce entirely new images or text with realistic characteristics

## Smooth Interpolation

Move smoothly between two latent points (e.g., morph one face into another seamlessly)

## Compact Representation

Latent vectors capture meaningful abstract features in low-dimensional space

## Controllable Generation

Manipulate  $z$  values to control specific aspects of generation (e.g., smiling → serious expressions)

# ◆ Mathematical Insight (Simplified)

The encoder learns to output two vectors:

- **Mean ( $\mu$ )** - central tendency
- **Standard deviation ( $\sigma$ )** - variability

A random latent vector is drawn as:

$$z = \mu + \sigma \times \varepsilon$$

where  $\varepsilon \sim \mathcal{N}(0, 1)$



## Reconstruction Loss

Measures how close the output is to the input

## KL Divergence Loss

Ensures latent space follows a normal distribution (smooth and continuous)

💡 Together, they make the model both **accurate** and **creative**.

## ◆ Component Summary

Component	Function	Analogy
Encoder	Learns to describe input as probabilistic features	Describes a face in numbers
Latent Space	Stores abstract, compressed representations	Sketchpad with "feature notes"
Decoder	Recreates or generates full data from latent vectors	Paints a new face from description



## ◆ Why VAEs Matter in Generative AI

- Learn *concepts*, not just copies  
Capture underlying patterns and structures in data
- Generate infinite new examples  
Create diverse, realistic samples from learned distributions
- Enable smooth transitions  
Interpolate between ideas seamlessly in latent space
- Foundation for modern models  
Basis for **Diffusion models** and **Text-to-Image generators**



### In One Sentence:

A Variational Autoencoder learns the **essence** of data as a probability distribution – allowing it to dream up new, realistic variations from its imagination (latent space).