

# Few-Shot Learning (FSL)

## How Machines Learn From Just a Few Examples


In this presentation, you'll discover what few-shot learning is, why it matters in modern AI, and how it differs from zero-shot and traditional learning approaches.

## The Big Idea

# "Learning a new concept with just a few examples — like humans do."

🧠 **Think about this:** You see just two photographs of a kiwi fruit and someone tells you what it is. The next time you spot a kiwi—whether in a market, on a menu, or in someone's kitchen—you recognise it instantly. That's precisely how few-shot learning works.

This mirrors human cognition, where we don't need hundreds of exposures to learn something new. A few well-chosen examples are often enough.



# Why Few-Shot Learning Matters



## Traditional ML Challenge

- Requires thousands of labelled examples
- Expensive data collection
- Struggles with rare classes



## Few-Shot Advantage

- Learns from just 1–10 examples per class
- Cost-effective and efficient
- Handles scarce data scenarios

💡 Few-shot learning is now being applied across **medicine** (detecting rare diseases), **robotics** (learning new object manipulation), **natural language processing** (understanding new languages), and **computer vision** (identifying rare species).



# Breaking Down the Terminology

Understanding the "shot" terminology helps clarify the spectrum of learning approaches:


Term	Meaning	Examples Needed
<b>Shot</b>	Number of labelled examples per class	Varies
<b>Zero-Shot</b>	No examples; uses semantic descriptions	0 examples
<b>One-Shot</b>	Learning from a single example	1 example
<b>Few-Shot</b>	Learning from a handful of examples	2–10 examples per class
<b>Many-Shot</b>	Traditional supervised learning	Hundreds to thousands

 **Few-shot** represents a practical middle ground—enough data to learn meaningful patterns, but not so much that data collection becomes prohibitively expensive.

# An Everyday Analogy



## Think Like a Chef



### Traditional ML Approach

Like a novice who needs to cook a dish hundreds of times, making mistakes and iterating constantly before mastering it.

### Few-Shot Learning Approach

Like an experienced chef who knows cooking fundamentals and can master a new dish after just 1–2 attempts with a recipe.



The key insight: The model leverages *prior knowledge* from related tasks to generalise quickly to new situations.

# Real-World Applications

Few-shot learning shines in domains where data scarcity is the norm, not the exception:

## Healthcare

### **Detecting rare diseases:**

Limited patient scans make traditional ML impractical. FSL can learn from just a handful of cases.

## Computer Vision

### **Identifying rare animals:**

When only a few photographs of endangered species exist, FSL enables accurate recognition.

## Robotics

### **Learning new object grasps:**

Physical demonstrations are expensive; robots can learn from just a few examples.

## Natural Language Processing

**New language tasks:** When minimal annotated text exists in low-resource languages, FSL bridges the gap.

# Comparing Learning Paradigms

Understanding where few-shot learning fits in the broader landscape of machine learning approaches:

Aspect	Zero-Shot	Few-Shot	Supervised
Examples Needed	0	1–10 per class	Hundreds or more
Knowledge Source	Semantic descriptions	Few labelled examples	Large labelled dataset
Best Use Case	Completely new unseen categories	Limited available data	Common, well-defined tasks
Example	"Zebra" learnt from text description only	"Kiwi" learnt from 3 images	"Cat" learnt from 1,000 images
Complexity	High	Medium	Low

# The Core Challenge

"How do we train a model that can learn to learn?"

## Generalise From Prior Experiences

The model must extract transferable knowledge from tasks it has seen before and apply that understanding to entirely new scenarios.

## Adapt to New Tasks Quickly

When presented with just a handful of examples, the model needs to rapidly adjust its internal representations without extensive retraining.

## Avoid Overfitting With Few Samples

With minimal data, there's a high risk of memorising specific examples rather than learning underlying patterns. The model must stay balanced.

# The Meta-Learning Paradigm

## Meta-Learning

### "Learning to Learn"

Instead of training a model to excel at one specific task, meta-learning trains a model to become good at *learning itself*. The model practises learning across many diverse small tasks, developing a generalised ability to adapt quickly.

Think of it as training someone not just to solve maths problems, but to become good at **learning new types of maths** efficiently.



# Key Takeaways

## FSL learns with very few labelled samples

Typically requiring just 1–10 examples per class, few-shot learning makes AI practical in data-scarce environments.

## Bridges the gap between zero-shot and supervised learning

FSL occupies the sweet spot—more practical than zero-shot, more efficient than traditional supervised approaches.

## Inspired by human-like adaptability

Just as humans can generalise from minimal exposure, FSL mimics our natural ability to recognise patterns quickly.

## Foundation for modern adaptive AI

From medical diagnosis to robotics, FSL is enabling the next generation of intelligent, flexible systems.

→ Coming up next: How Few-Shot Learning Works – Techniques & Architectures