

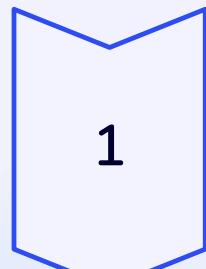
How Few-Shot Learning Works

Meta-Learning, Models & Real-World Use

A comprehensive guide to understanding how machines learn from limited data, inspired by human cognitive abilities.

Core Process Overview

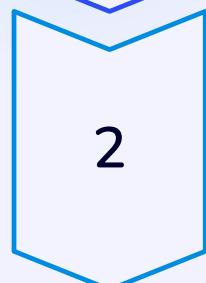
Few-shot learning follows a systematic approach that enables models to generalise from minimal examples. The process mirrors human learning patterns, where we quickly adapt to new concepts after seeing just a few instances.



Training Phase: Learn to Learn

1

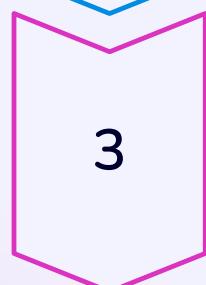
The model is exposed to diverse mini-tasks, developing meta-knowledge about how to extract meaningful patterns efficiently.



Inference Phase: Adapt to New Classes

2

Using acquired meta-knowledge, the model quickly adapts to classify previously unseen categories with just a few labelled samples.



Prediction: Classify Accurately

3

The adapted model confidently predicts labels for new examples, demonstrating true generalisation capability beyond rote memorisation.

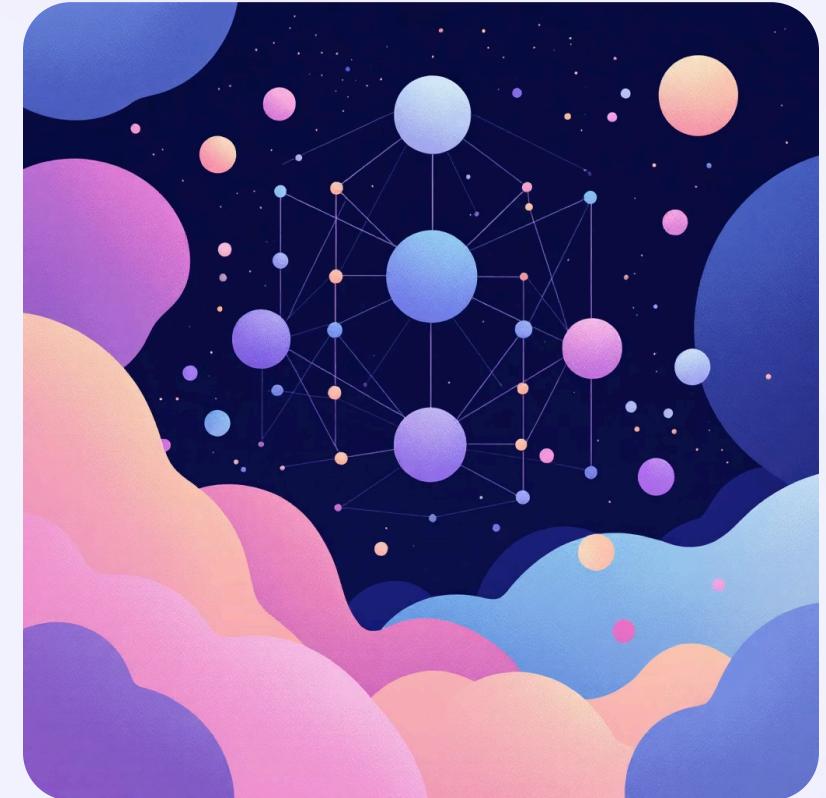
Phase 1: Training (Meta-Learning)

During the training phase, the model learns generalisation strategies rather than memorising specific patterns. This is achieved through exposure to numerous mini-tasks.

Each Mini-Task Contains:

- **Support set:** A few labelled samples per class (typically 1-5 examples) that serve as reference points
- **Query set:** Unlabelled examples the model must classify using only the support set

This process is repeated across hundreds or thousands of randomly generated task combinations, forcing the model to develop robust feature extraction and comparison mechanisms.



The key insight: By solving diverse tasks, the model learns *how to learn* rather than what to learn.

Example of Meta-Tasks

Consider how the model is trained across varied classification tasks. Each task presents a unique challenge, yet shares underlying patterns that the model must discover.

Task	Classes	Support Set	Query Set
Task 1	Cats vs Dogs	5 examples each	10 unlabelled
Task 2	Apples vs Oranges	5 examples each	10 unlabelled
Task 3	Sedan vs SUV	5 examples each	10 unlabelled

What the Model Learns

Through this diverse training regime, the model develops an understanding of **shared patterns** across different domains—shapes, textures, semantic relationships, and discriminative features that generalise beyond specific categories.



Key Architectures

1. Prototypical Networks

Prototypical Networks create a compact representation for each class by computing a **prototype**—essentially the average embedding of all support examples. Classification then becomes a simple distance measurement problem.

Mathematical Foundation:

$$c_k = \frac{1}{N_k} \sum_{i=1}^{N_k} f(x_i)$$

Where:

- c_k : prototype (centroid) for class k in embedding space
- $f(x_i)$: feature embedding of support sample i , extracted via a learnt encoder
- N_k : total number of support samples in class k

□ Query examples are classified by finding the nearest prototype using Euclidean distance—the class whose prototype is closest wins. This elegant approach reduces few-shot learning to a metric learning problem.

Key Architectures (continued)

2. MAML

Model-Agnostic Meta-Learning

MAML learns initialisation parameters that serve as an excellent starting point for rapid adaptation. After just a few gradient descent steps on new task data, the model achieves strong performance.

Key advantage: Works with any gradient-based model—CNNs, RNNs, or Transformers.

3. Matching Networks

Attention-Based Similarity

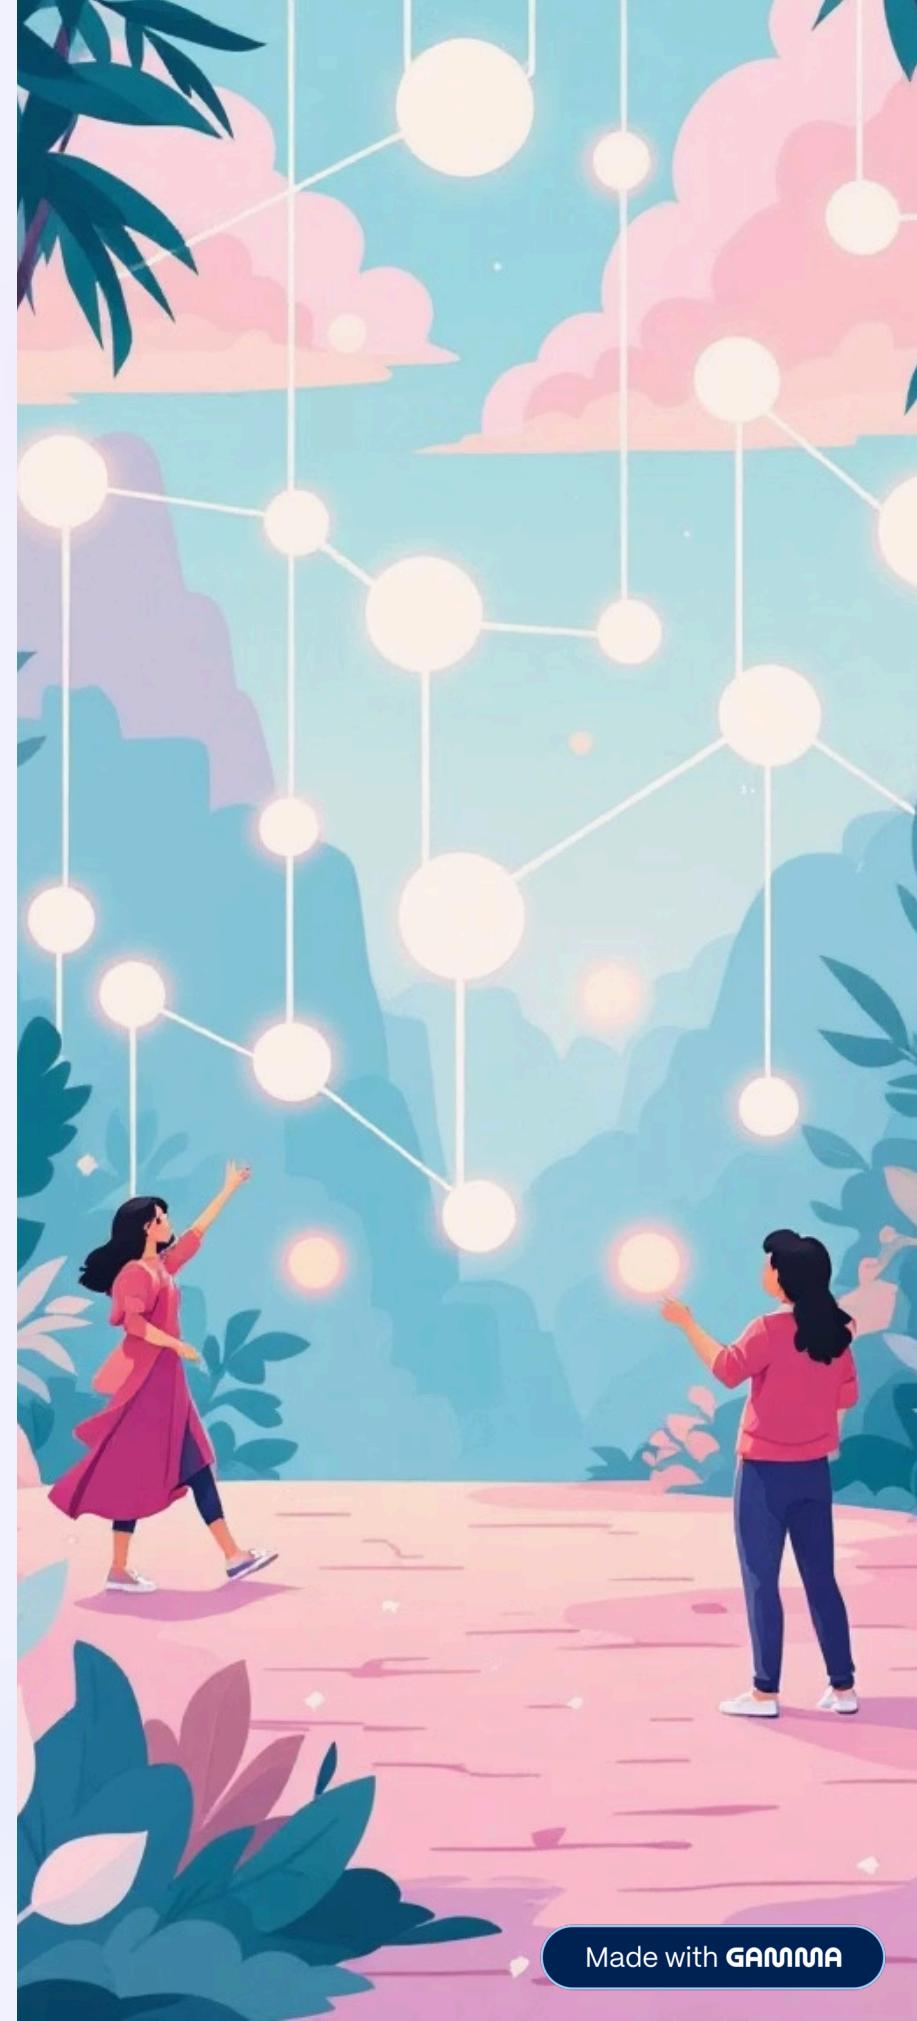
Rather than creating prototypes, Matching Networks directly compare new inputs to individual support examples using attention mechanisms and similarity metrics.

Classification uses weighted voting based on similarity scores.

Cosine Similarity in Matching Networks:

$$\text{cosine}(x, y) = \frac{x \cdot y}{\|x\| \|y\|}$$

This metric measures the angle between feature vectors, emphasising directional similarity whilst being invariant to magnitude—ideal for comparing embeddings of different scales.



Phase 2: Inference (New Task)

Now let's see few-shot learning in action with a practical example. Suppose we need to build a fruit classifier but have limited training data.

Given Support Set:

- 3 images of kiwi 🍁
- 3 images of mango 🥭

Goal:

Classify a new, unseen fruit image into one of these two categories.



01

Extract Features

Pass all support images through the trained feature extractor network to obtain high-dimensional embeddings.

02

Build Prototypes

Compute class prototypes by averaging the embeddings—one prototype for kiwi, one for mango.

03

Classify Query Image

Extract features from the new image, measure distances to both prototypes, and assign the label of the nearest prototype.

- ❑ The entire inference process requires no gradient updates or retraining—just forward passes through the network. This enables real-time deployment in resource-constrained environments.

Modern Approaches

Contemporary few-shot learning leverages the power of large-scale pre-trained models, which have already learnt rich representations from massive datasets. This transfer learning paradigm significantly improves few-shot performance.

Pre-trained Language Models

BERT, GPT, T5

These models encode deep linguistic knowledge from billions of text tokens. With prompt engineering or minimal fine-tuning on 2-10 examples, they achieve impressive few-shot text classification, sentiment analysis, and entity recognition.

Pre-trained Vision Models

ResNet, Vision Transformers, CLIP

Trained on ImageNet or web-scale image-text pairs, these encoders capture universal visual features. Few-shot adaptation enables recognition of novel object categories, medical imaging tasks, and satellite image classification.

Practical Example:

GPT-3 can classify customer reviews as positive/negative after seeing just 2 labelled examples in the prompt – no gradient updates required. This is known as **in-context learning**.

Advantages & Challenges

Like any machine learning paradigm, few-shot learning presents compelling benefits alongside significant technical hurdles. Understanding both is crucial for successful deployment.

Advantages

- **Data efficiency:** Works effectively with minimal labelled data, reducing annotation costs dramatically
- **Rapid adaptation:** Quickly learns new tasks without extensive retraining, enabling agile deployment
- **Human-like learning:** Mimics human cognitive ability to generalise from few examples
- **Practical applicability:** Invaluable in domains where data collection is expensive, dangerous, or ethically constrained

Challenges

- **Overfitting risk:** Models can memorise few examples rather than learning robust features
- **Dependency on meta-training:** Requires large, diverse meta-training datasets to develop generalisation skills
- **Computational intensity:** Meta-learning involves training across thousands of tasks—resource-intensive
- **Task similarity:** Performance degrades significantly when test tasks differ substantially from meta-training distribution



Summary



Human-Style Learning

Few-shot learning enables AI systems to generalise from limited examples, mirroring human cognitive efficiency.



Powered by Meta-Learning

Through meta-learning and transfer from pre-trained models, systems develop the ability to "learn how to learn."



Proven Architectures

Prototypical Networks, MAML, and Matching Networks provide robust frameworks for few-shot classification.



Real-World Impact

Applications span healthcare diagnostics, robotics, NLP, computer vision, drug discovery, and personalised recommendations.

"Few-Shot Learning teaches AI to learn—not just to remember."

As data annotation costs rise and privacy concerns grow, few-shot learning represents a critical pathway towards more efficient, adaptable, and human-like artificial intelligence systems.